Weekly Report(May 21 - May 27)

Liu Junnan

Abstract

This week I finished cs231n course and started to do assignment3.

1 Work Done

1.1 RNN and LSTM

RNN Unlike feed forward neural networks, recurrent neural networks take as their input not just the current input example they see, but also what they have perceived previously in time. The decision a recurrent net reached at time step $t-1$ affects the decision it will reach one moment later at time step t. So RNNs have two sources of input, the present and the recent past. The sequential information is preserved in the recurrent network's hidden state, which manages to span many time steps as it cascades forward to affect the processing of each new example. One way to think about RNNs is this: they are a way to share weights over time. Therefore RNNs are good at processing sequences like image caption and machine translation.


Figure 1: RNN unit

The forward pass of a RNN block can be mathematically described as:

$$h_t = f_W(h_{t-1}, x_t) = W_h^T \cdot h_{t-1} + W_x^T \cdot x_t \tag{1}$$

where h_t is the current hidden state at time step t; f_W the function – either sigmoid or tanh – that squashes the sum of the weight input and hidden state, making gradients workable for backpropagation; h_{t-1} the previous hidden state; x_t the input at the same time step.

Recurrent networks rely on an extension of backpropagation called backpropagation through time, or BPTT. Time, in this case, is simply expressed by a well-defined, ordered series of calculations linking one time step to the next, which is all backpropagation needs to work.

However, RNNs suffer a lot from vanishing/exploding gradient problems. Intuitively, backpropagation from h_t to h_{t-1} multiplies by W_h, so computing gradient of h_0 involves many factors of W_h. If the largest singular value of W_h is greater than 1, then exploding gradients would occur; on the contrary if the largest singular value of W_h if less than 1, then vanishing gradients would happen.

1

Exploding gradients can be solved relatively easily, because they can be truncated or squashed. Vanishing gradients can become too small for computers to work with or for networks to learn – a harder problem to solve.

LSTM   Long short-term memory was proposed to solve vanishing gradient problem of vanilla RNNs. LSTMs design complicated cells that help preserve the error that can be backpropagated through time and layers.
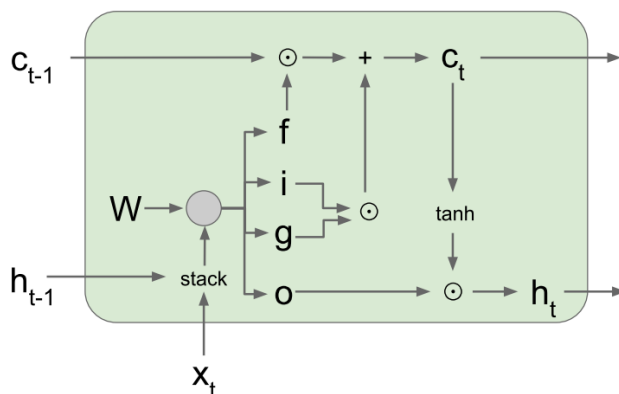

Figure 2: LSTM cell

The forward pass of a LSTM cell is defined as follows:

$$\begin{pmatrix} i \\ f \\ o \\ g \end{pmatrix} = \begin{pmatrix} \sigma \\ \sigma \\ \sigma \\ \tanh \end{pmatrix} W \begin{pmatrix} h_{t_1} \\ x_t \end{pmatrix} \tag{2}$$

$$c_t = f \odot c_{t_1} + i \odot g$$
$$h_t = o \odot \tanh(c_t)$$

where $\odot$ is element-wise multiplication.

As described in equation(2), all the computations involved, including sigmoid, element-wise multiplication and so on, are differentiable, which are suitable for backpropagation.

Those gates act on the signals they receive, and similar to the neural network's nodes, they block or pass on information based on its strength and input, which they filter with their own sets of weights. Those weights, like the weights that modulate input and hidden states, are adjusted via the recurrent networks learning process. That is, the cells learn when to allow data to enter, leave or be deleted through the iterative process of making guesses, backpropagating error, and adjusting weights via gradient descent.

1.2   Generative Models

Neural network models we have learned so far, such as feed forward networks, CNNs, RNNs, LSTMs, are all basically discriminative models, which directly estimate posterior probabilities. Generative models focus on modeling class-conditional probabilistic distribution functions and prior probabilities, and a good outcome is that generative models can generate synthetic data points. Therefore generative adversarial networks become an active research field last year due to this feature.

CS231 course discussed generative models like PixelRNN, PixelCNN, variational autoencoder and GAN, and used them to generate instances or do other interesting tasks. Although the theoretical derivations are difficult to understand, but perhaps the assignment will offer practical approach to better understanding.

2

1.3 RNN - Implementation

The forward pass of RNN is defined by equation(1), and the code is simply one line:

```
1  next_h = np.tanh(np.dot(x, Wx) + np.dot(prev_h, Wh) + b)
```

Since all the computations are differentiable, the backward pass is also simple:

```
1  dh = (1 - next_h**2) * dnext_h
2  dx = np.dot(dh, Wx.T)
3  dprev_h = np.dot(dh, Wh.T)
4  dWx = np.dot(x.T, dh)
5  dWh = np.dot(prev_h.T, dh)
6  db = np.sum(dh, axis=0)
```

Above are what a RNN unit does in a single time step, but usually a RNN unit will process multiple time steps. At each iteration we should assign next_h of last iteration to prev_h of current iteration.

1.4 LSTM - Implementation

The implementation of LSTM uses a big matrix computed by weights W, hidden states h and input x, and then splits it to i, f, o, g gates.

```
1  a = np.dot(x, Wx) + np.dot(prev_h, Wh) + b
2  ai, af, ao, ag = np.split(a, 4, axis=1)
3
4  i = sigmoid(ai)
5  f = sigmoid(af)
6  o = sigmoid(ao)
7  g = np.tanh(ag)
8
9  next_c = f * prev_c + i * g
10 next_h = o * np.tanh(next_c)
```

However, there is a tricky part in backward pass. Notice that c_t also shows up in the equation that computes h_t, so ∇c_t backpropagated from next cell should also plus $\frac{\partial h_t}{\partial c_t}$.

2 Plans

In the next I plan to finish assignment 3 of cs231 and review generative model part of the course.